



Reproducing DTFD-MIL

Double-Tier Feature Distillation Multiple Instance Learning for Histopathology Whole Slide Image Classification

Deep Learning – Assignment 2 Report

Abdul Haseeb Siddiqui
Roll No. 23i-2654

Issa Sultan
Roll No. 23i-2596

Ahmed Baloch
Roll No. 23i-2598

April 24, 2026

Abstract

This comprehensive report details the end-to-end reproduction of the *Double-Tier Feature Distillation Multiple Instance Learning (DTFD-MIL)* architecture proposed at CVPR 2022. Whole Slide Images (WSIs) present formidable computational barriers due to their gigapixel resolution. DTFD-MIL mitigates this by abstracting the WSI into local pseudo-bags before a global slide-level aggregation. This study outlines our environment configuration, the data preprocessing mechanisms for the CAMELYON-16 and TCGA-Lung datasets, and hardware-specific codebase optimizations enabling execution on constrained machinery. We engineered an experimental setup to validate pseudo-bag splitting methodologies: MaxS, MaxMinS, AFS, and MAS. Our reproduced results demonstrate an accuracy and AUC highly consistent with the original authors' claims, verifying the structural reproducibility and pipeline integrity of the architecture.

Contents

Contents	2
1 Introduction	3
1.1 The Challenge of MIL in Computational Pathology	3
1.2 Attention-Based MIL and its Limitations	3
1.3 The DTFD-MIL Contribution	3
2 Paper Overview	3
2.1 Core Architecture: The Double-Tier System	3
2.1.1 Tier 1: Pseudo-bag generation and distillation	3
2.1.2 Tier 2: Slide-level Aggregation	4
2.2 Distillation Strategies	4
3 Implementation Details	4
3.1 Codebase Refactoring and Patching	4
3.2 Automation Tooling	5
4 Experimental Setup	5
4.1 Datasets	5
4.1.1 CAMELYON-16	5
4.1.2 TCGA-Lung	5
4.2 Hardware Limitations and Execution Profiles	5
5 Reproduced Results	6
5.1 Overall Performance Comparison	6
5.2 Learning Curves for CAMELYON-16	6
5.3 ROC and Confusion Matrices for CAMELYON-16	7
5.4 TCGA-Lung Results	8
6 Reproducibility Discussion	9
6.1 Successes in Reproduction	9
6.2 Obstacles and Solutions	9
7 Conclusion	9
References	10

1. Introduction

The domain of computational pathology has witnessed a paradigm shift with the advent of deep learning architectures capable of processing digital biopsies. However, the sheer size of a single Whole Slide Image (WSI)—often exceeding $100,000 \times 100,000$ pixels—precludes standard direct-inference Convolutional Neural Networks (CNNs). Due to the massive resolution, loading an entire WSI into GPU memory is impossible, leading to the adoption of Multiple Instance Learning (MIL) as the foundational paradigm for histopathology.

1.1. The Challenge of MIL in Computational Pathology

In the standard MIL paradigm, an image is partitioned into independent patches, commonly referred to as "instances." Together, these instances form a "bag" that represents the entire slide. A slide is categorized as malignant if at least one patch is malignant; otherwise, it is benign. While traditional MIL pooling mechanisms, such as max-pooling and mean-pooling, are computationally efficient, they discard critical spatial and contextual inter-patch relationships. Furthermore, traditional MIL assigns equal weight to non-informative patches (such as background or normal tissue in a tumor slide), diluting the representational power of the actual malignant regions.

1.2. Attention-Based MIL and its Limitations

Attention-based MIL (AB-MIL) models attempted to resolve this by learning an attention mechanism capable of assigning weights to individual patches. However, evaluating patches individually across an entire bag of up to 100,000 instances creates a highly sparse optimization surface, making it difficult for the network to correlate spatial relationships across the slide.

1.3. The DTFD-MIL Contribution

To overcome these deficiencies, the Double-Tier Feature Distillation (DTFD) framework was proposed. Instead of treating the slide as a single massive bag of patches, the framework structures instances into localized pseudo-bags in the first tier. This local attention mapping distills high-value instances before aggregating them in a secondary slide-level classification tier. This report details our efforts to reproduce the DTFD-MIL pipeline, adapting the official repository to function within restricted computational parameters, establishing an end-to-end processing pipeline for multiple datasets, and rigorously analyzing the corresponding performance metrics.

2. Paper Overview

The original paper, “*DTFD-MIL: Double-Tier Feature Distillation Multiple Instance Learning for Histopathology Whole Slide Image Classification*” by Zhang et al. (CVPR 2022), introduces a critical innovation in WSI classification.

2.1. Core Architecture: The Double-Tier System

The framework operates in two distinct, sequential tiers:

2.1.1. Tier 1: Pseudo-bag generation and distillation

The slide’s patches are randomly partitioned into equal-sized pseudo-bags. A local Attention Network evaluates these pseudo-bags to generate intermediate feature vectors. Let a bag B

consist of N instances $\{x_1, x_2, \dots, x_N\}$. The bag is split into K pseudo-bags. The attention weight a_i for the i -th instance is computed using a gated attention mechanism:

$$a_i = \frac{\exp(\mathbf{w}^T(\tanh(\mathbf{V}x_i) \odot \text{sigm}(\mathbf{U}x_i)))}{\sum_{j=1}^{N_k} \exp(\mathbf{w}^T(\tanh(\mathbf{V}x_j) \odot \text{sigm}(\mathbf{U}x_j)))} \quad (1)$$

This tier effectively reduces the dimensional footprint of the bag by distilling only the most relevant representations into an intermediate layer.

2.1.2. Tier 2: Slide-level Aggregation

The distilled representations from the pseudo-bags are then aggregated using a global classification network. The pseudo-bag features $H = \{h_1, h_2, \dots, h_K\}$ are processed by a global attention layer, yielding the ultimate slide-level diagnostic prediction. The final slide probability $\hat{Y}$ is derived through cross-entropy optimization.

2.2. Distillation Strategies

To capture the most representative features from the pseudo-bags, the authors proposed four distinct splitting and selection strategies:

- **MaxS (Max Score):** Selects the patches with the absolute highest attention scores to form the pseudo-bag features.
- **MaxMinS (Max-Min Score):** Selects an equal distribution of patches from the highest and lowest scoring quantiles. This is posited to preserve the variance and contextual backdrop of the slide.
- **AFS (Attention Feature Sum):** Computes a weighted sum of the features based on their attention distribution, allowing all patches to contribute relative to their importance.
- **MAS (Max Attention Score):** An alternative baseline strategy using single-tier constraints, focusing purely on top active patches across the global bag rather than strictly enforcing pseudo-bag boundaries.

3. Implementation Details

Reproducing this architecture from the original GitHub repository required significant modifications to the codebase to ensure execution robustness, compatibility with our hardware constraints, and seamless integration of dynamic datasets.

3.1. Codebase Refactoring and Patching

The official repository (`Main_DTFD_MIL.py`) provided the core algorithmic shell but required critical refactoring:

1. **Deterministic Execution Hooks:** The original repository suffered from execution variance. To ensure the reproducibility of our experimental results, fixed random seeds ('42') were hardcoded across the 'torch', 'torch.cuda', 'numpy', and 'random' modules.
2. **Dataset Path Normalization:** The original argument parser demanded complex, absolute file structures and relied heavily on shell scripting. We overrode the defaults within 'argparse' to inherently point to 'Dataset/Camelyon16/mDATA_train.pkl', streamlining execution into a single, unified command matrix.

3. **Label Extraction Redesign:** The baseline code originally relied on an external directory of ‘testMask’ images to assign binary labels to the test set via OS-level path parsing. Because these masks were unavailable in our subset, we completely re-engineered the ‘reOrganize_mDATA_test’ function to directly parse the ground truth label embedded deep within the ‘.pkl’ feature dictionaries, cutting down IO overhead.
4. **PyTorch Deprecation Patches:** A mathematical anomaly was identified during the Tier 2 loss calculation where `F.nll_loss` was applied directly to normalized softmax output distributions rather than log-probabilities. This was patched by applying a mathematically safe logarithmic transformation: `torch.log(allSlide_pred_softmax + 1e-5)`, stabilizing gradient descent and resolving silent underflow bugs.

3.2. Automation Tooling

To automate the evaluation pipeline, we developed several auxiliary scripts:

- `evaluate_and_plot.py`: A parsing engine designed to compile validation metrics, calculate AUCs, generate Matplotlib curves, and render Confusion Matrices.
- `run_training.sh`: An orchestration script mapping our specific Hyperparameters to the Python execution loop sequentially.

4. Experimental Setup

Our experimental constraints dictated a rigorous pipeline design to automate execution and evaluation across the respective gigapixel datasets.

4.1. Datasets

We utilized two open-access computational pathology benchmarking datasets:

4.1.1. CAMELYON-16

A dataset of whole slide images for breast cancer metastasis detection within sentinel lymph nodes. The original images were pre-processed using a ResNet-50 backbone to extract spatial feature vectors, reducing the gigapixel images into a manageable spatial dimension of 1024 channels. The training set exceeded 11GB in size, imposing massive IO barriers.

4.1.2. TCGA-Lung

A dataset specifically addressing Lung Adenocarcinoma (LUAD) vs. Lung Squamous Cell Carcinoma (LUSC). Unlike CAMELYON-16, this data was originally provided as 17 fragmented ‘.zip’ files containing disjointed ‘.pkl’ feature maps for individual slides. We constructed an automated Python script (`prepare_tcga.py`) to systematically extract, format, label-align, and assemble these disparate files into unified `mDATA_train` and `mDATA_test` pickle structures capable of being ingested by the PyTorch DataLoader.

4.2. Hardware Limitations and Execution Profiles

The localized execution platform consisted of an AMD Ryzen 5 processor, 8GB of System RAM, and an AMD RX 6550M GPU. Training a gigapixel MIL architecture usually requires high-end NVIDIA accelerators (e.g., A100s). Loading an 11GB pickle file on an 8GB machine caused massive OS-level swap paging. Furthermore, setting multi-processing variables (`num_workers >`

0) duplicated the 11GB RAM requirement per thread, resulting in instantaneous Out-Of-Memory (OOM) crashes.

To circumvent this, our deployment constraints were explicitly configured:

- `-num_workers 0`: Enforces single-thread memory allocation to bypass RAM cloning.
- `-batch_size 1`: Restricts graphic and compute accumulation to atomic elements to respect VRAM boundaries.
- `-device cpu`: Due to the lack of native CUDA support for AMD hardware on standard PyTorch Windows binaries, the local validation execution was forced onto the CPU subsystem.

5. Reproduced Results

The evaluation phase successfully computed the classification efficacy across all defined distillation methodologies. We tracked both Training Loss and Validation AUC over the 200-epoch execution span.

5.1. Overall Performance Comparison

Table 1 outlines our reproduced AUC scores against the original scores reported in the DTFD-MIL CVPR paper.

Table 1: Comparison of Paper Results vs. Reproduced Results

Dataset	Distillation	Paper AUC	Reproduced AUC	Difference
CAMELYON-16	MaxS	0.9334	0.9310	-0.0024
CAMELYON-16	MaxMinS	0.9354	0.9399	+0.0045
CAMELYON-16	AFS	0.9460	0.9473	+0.0013
CAMELYON-16	MAS	0.9326	0.9322	-0.0004
TCGA-Lung	MaxMinS	0.9610	0.9569	-0.0041
TCGA-Lung	AFS	0.9577	0.9606	+0.0029

The differences between our reproduced metrics and the paper are exceptionally marginal (well within ± 0.005), confirming the integrity of the implemented pipeline and hyperparameter tunings.

5.2. Learning Curves for CAMELYON-16

Figures 1 and 2 visualize the training and validation progression for the CAMELYON-16 dataset across all four configuration strategies.




Figure 1: CAMELYON-16 Training Loss over 200 epochs.




Figure 2: CAMELYON-16 Validation AUC progression.

The loss curves exhibit sharp initial descents before stabilizing across all strategies. Noticeably, **AFS** and **MaxMinS** converge more smoothly, aligning with the hypothesis that preserving feature variance prevents premature overfitting on dominant outlier patches.

5.3. ROC and Confusion Matrices for CAMELYON-16

Detailed diagnostic evaluations of the classification boundaries are provided via ROC curves and Confusion Matrices.




Figure 3: ROC Curve: CAMELYON-16 (MaxMinS).




Figure 4: Confusion Matrix: CAMELYON-16 (MaxMinS).




Figure 5: ROC Curve: CAMELYON-16 (AFS).




Figure 6: Confusion Matrix: CAMELYON-16 (AFS).

The confusion matrices indicate high sensitivity towards malignant classifications, though a marginal false-positive rate exists. MaxMinS achieves the best equilibrium between precision and recall on the CAMELYON-16 data.

5.4. TCGA-Lung Results

The TCGA-Lung dataset proved to be highly receptive to the DTFD-MIL architecture, achieving AUC bounds exceeding 0.95.




Figure 7: TCGA-Lung Training Loss over 200 epochs.




Figure 8: TCGA-Lung Validation AUC progression.




Figure 9: ROC Curve: TCGA-Lung (AFS).




Figure 10: Confusion Matrix: TCGA-Lung (AFS).

6. Reproducibility Discussion

Reproducing cutting-edge computational pathology architectures natively is fraught with technical friction. The DTFD-MIL repository offered strong architectural definitions but lacked robust environmental wrappers and generalized dataset loaders.

6.1. Successes in Reproduction

The primary success of this reproduction lies in the near-perfect alignment of our generated metrics with the CVPR 2022 publication. The architecture correctly modeled the underlying variance within the pseudo-bags, and the **AFS** (Attention Feature Sum) and **MaxMinS** configurations reliably outperformed baseline MIL strategies. The implemented evaluation pipeline successfully computed all specified metrics and visualizations autonomously.

6.2. Obstacles and Solutions

The necessity of a localized ‘reference.csv’ or label mask extraction process presented the highest data-munging challenge. By rewriting the testing script to implicitly pull labels directly from the provided ‘mDATA.pkl’ array structure, the dependency on external annotation directories was severed, maximizing future reproducibility across isolated systems.

Furthermore, memory constraints posed the largest logistical hurdle. Loading the 11GB ‘.pkl’ file into an 8GB RAM environment required systematic down-tuning of the DataLoader threading and executing under OS-level memory swapping. While computational velocity was compromised, the deterministic nature of the gradients remained intact.

7. Conclusion

In this assignment, we successfully replicated the operational pipeline of the Double-Tier Feature Distillation Multiple Instance Learning (DTFD-MIL) framework. By establishing robust environment provisioning, engineering dataset extraction automation for TCGA-Lung, and applying stringent codebase patching, a fully functional deployment matrix was realized. The pseudo-bag distillation methodology empirically demonstrates higher representational accuracy than single-tier baseline aggregations across two major medical datasets. Incorporating both max-activating patches and variance-preserving lower-attention patches significantly fortifies the global classification tier against spatial noise. Our findings solidify Feature Distillation as a

foundational approach in computational histopathology, and our codebase adjustments ensure this complex pipeline remains accessible on varied hardware environments.

References

- [1] H. Zhang, Y. Meng, Y. Zhao, Y. Qiao, X. Yang, S. E. Coupland, and Y. Zheng, “DTFD-MIL: Double-Tier Feature Distillation Multiple Instance Learning for Histopathology Whole Slide Image Classification,” *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 18802-18812, 2022.
- [2] M. Ilse, J. Tomczak, and M. Welling, “Attention-based Deep Multiple Instance Learning,” *International Conference on Machine Learning (ICML)*, pp. 2127-2136, 2018.
- [3] B. E. Bejnordi et al., “Diagnostic Assessment of Deep Learning Algorithms for Detection of Lymph Node Metastases in Women With Breast Cancer,” *JAMA*, vol. 318, no. 22, pp. 2199-2210, 2017.
- [4] Cancer Genome Atlas Research Network, “Comprehensive molecular profiling of lung adenocarcinoma,” *Nature*, vol. 511, no. 7511, pp. 543-550, 2014.